

# EEE3096S: Embedded Systems II

## LECTURE 4: BENCHMARKING

Presented by:  
**Dr Yaaseen Martin**



Electrical Engineering  
University of Cape Town

# SELECTION PROCESS

- **Selection Process:**  
Selecting parts and other things from which you will build your embedded system
- This includes choosing processor(s), hardware components, software and code to incorporate and/or reuse.



# THE SELECTION PROCESS

- A significant part of the embedded design process
- Involves selecting the best software and hardware choices for the required **functionality**
- Correct selection is **critical to profit \$\$\$**
- Can (greatly) lower production costs
  - No money wasted on unused features
- Lower non-recurring engineering costs
  - Development cheaper with good tools
- Better reliability and maintainability



## THE 'GOLDEN MEASURE'



## TERM: GOLDEN MEASURE

- This concept originates from the theory of a ‘gold standard measure’\* which is:
  - a benchmark available under reasonable conditions. Not necessarily the perfect test, but the best available one that has a standard with known results. This is especially important when faced with the impossibility of direct measurements
- A **golden measure** in computing could be:
  - A solution that may run slowly, isn’t optimized, but *you know* it gives (numerically speaking) excellent results

e.g., a solution written in MATLAB or Python, verify it is correct using graphs or exhaustive testing, or checking by hand with calculator, etc.

\* This term originated more in medicine (although has a much older origin if looking further back) and is also widely applied to data science and computing nowadays.

# BENCHMARKING

- **Benchmarking:**  
Process of measuring performance of a system (an embedded system in our case)
- It's a program (or systematic approach) that **quantitatively evaluates** performance, cost, and computing resources (among other things) of a computing solution
- **Benchmark suite:**  
A sets of benchmark programs designed to get a comprehensive view of the performance of a computer system for executing a variety of representative processing operations.

# WALL CLOCK TIME

- Generally the most “accurate”: use a built-in timer, which is directly related to **real time** (e.g., if the timer measures 1s, then 1s elapsed in the real world)
- Note: To avoid overflow, used unsigned variables

```
start = read_the_timer();  
    DO PROCESSING  
end = read_the_timer();  
time = end-start
```



# STDC: GETTIMEOFDAY

## **gettimeofday**

- Very portable, part of the StdC (Standard C) library
- Should be available on any Linux system
- Returns time in seconds and microseconds since midnight January 1 1970
- Uses struct timeval comprising
  - tv\_sec : number of seconds
  - tv\_usec : number of microseconds\*
- Converting to microseconds will use huge numbers, rather work on differences

**\* Word of caution: some implementations always return 0 for the usec field!!**

*On Cygwin, the resolution is only in milliseconds, so tv\_usec in multiples of 1000. Not provided in DevC++ on Windows.*



# GETTIMEOFDAY EXAMPLE

```
#include <stdio.h>
#include <sys/time.h>
#include <time.h>
struct timeval start_time, end_time; // variables to hold start and end time

int main() {
    int tot_usecs; /* may want to use long int on CPU with smaller word length */
    gettimeofday(&start_time, (struct timezone*)0); // starting timestamp

    /* Do some work */
    int i,j,sum=0;
    for (i=0; i<10000; i++)
        for (j=0; j<i; j++) sum += i*j;

    gettimeofday(&end_time, (struct timezone*)0); // ending timestamp
    tot_usecs = (end_time.tv_sec-start_time.tv_sec) * 1000000 +
        (end_time.tv_usec-start_time.tv_usec);
    printf("Total time: %d usec.\n", tot_usecs);
}
```

# WHAT IS WRONG ABOUT USING ONLY WALL CLOCK TIME?

It can provide a false impression of how effective your solution is – or rather more correctly it at least **doesn't give a 'full picture' of performance** ...

Things to consider:

- Typically do tests after the system has 'warmed up' (cache loaded) by running the same benchmark multiple times\*
- Speed improved but accuracy sacrificed?
- Development effort vs. execution speed improvement?
- Resource costs for upgrading vs. costs saved by remaining with the old version?
- Power usage? Does the new solution need more power (per execution, also on average including idle time)
- Maintainability? (e.g. is the new version more complex?)
- Environment impact? (Does the upgrade result in waste that could be environmentally detrimental)

\* But note, this could also be a false impression depending how your system is likely to work, if operations are typically not going to be in cache when invoked, this approach could invalidate your benchmarking.

# BENCHMARKING: WHAT TO TEST

- Compiler
  - Converts High Level Language to Assembly Language thus we benchmark compiler efficiency (such as how efficient is the generated assembly code?)
- Processor
  - Code in hand-crafted/inspected assembly (to make comparisons fair)
- Operating System
  - Interrupt latencies, overhead of operating system calls, limits on devices, kernel size, availability of services and facilities such as support for virtual memory and paged memory.
- Platform
  - Scalability of memory. Peripheral limits. Interfaces supported. Power use. Power saving features. OS's supported.
- Applications

## BENCHMARKING: WHAT IS USUALLY MEASURED

- For embedded and DSP systems we typically benchmark:
  - Cycle count
  - Data and Program memory usage
  - Execution time
  - Power consumption (has recently become a common thing to report on)

# PROCESSOR PERFORMANCE

- This is usually talked about the most, and is often the most important design decision
- *Particularly:* will the proposed processor/controller/DSP be fast enough?
- Performance measuring tools are available to analyse this, at least roughly, particularly:
  - MIPS (Million Instructions Per Second)
  - Dhrystone Benchmark
  - FLOPS (Floating Point Operations Per Second)


**An important point to note:** that processor manufacturer might optimise their toolchains to get better benchmark results using these common tools, so it's usually best to use some more representative tests for your particular application.



# EXAMPLE OF STANDARD BENCHMARK

- The Dhrystone benchmark
  - Developed in 1984 by R. Wecker
  - Tests integer performance only
    - No I/O, no OS calls, no floating point
- Good for estimating integer performance
- Dhrystone score is the number of times a small program can run per second
- Like all benchmarks, applicability of Dhrystone varies (may be totally irrelevant to your specific application)

*... so don't get too comfortable using a common very 'general purpose' benchmarking tool like Dhrystone!*





# MEANINGFUL BENCHMARKING

- How will the processor perform with our application?
- Including complex factors
  - Interrupt handling
  - Task switching
  - Memory performance
- Even more important for Real Time applications
  - Context switching time is critical
  - Save registers, change mode, etc.

# RUNNING CUSTOM BENCHMARKS

- Typically use evaluation boards
  - Save costs of designing support hardware
- Evaluation boards
  - Single board computers
  - Include a selection of common peripherals
- Some manufacturers sell **Hot Boards**
  - Designed to boost benchmarks
  - Watch out for these boards



## BAD NEWS ABOUT BENCHMARKS

- Benchmarks do not consider all design factors which have an effect on performance, such as:
  - Power consumption
  - Physical size
  - Board layout
- Do not rely entirely on benchmark results for processor selection

# WHAT IS A CODE REVIEW?

- It is a practice (e.g. meeting or regular team-link up) of reviewing code developed for a project
- It follows (generally) a systematic practice  
*(formulate your own initiative method for now!)*
- It's about managing quality of code and ways to assess this, e.g. consider:
  - How to spot and log any defects
  - How many or common are the defects
  - Methods to use, such as (not necessarily all of these, not necessarily used in every review) are:
    - Unit testing
    - Function testing
    - Integration testing
    - How can this be improved?